

Liveraging Functional Reactive Programming to Build Modern NFV Applications

Abstract

Abstract

1 Introduction

In recent years, the research community has witnessed the quick development of Network Function Virtualization (NFV). NFV aims to move away from hardware-based network functions (NFs) and run software NFs in virtualized environment such as Virtual Machines (VMs) and containers. NFV greatly simplifies the deployment of NFs on critical network paths, as it no longer requires the network operators to maintain proprietary hardware NFs.

Over the past years, the performance, scalability and programmability of NFV have all been greatly improved. However, the core of NFV still lies in how to efficiently implement NF software. Even though there are many research works [] discussing advancements in implementing NF software, these works are all done in an ad-hoc manner, i.e. the advancement of these works are made using a wide range of tools and methods, lacking a systematic framework to capture all of these advancements.

On the other hand, there are a wide range of NFs that must be implemented for NFV, such as firewall, NAT, asset monitor and IDS. These NFs operate on different layers, i.e. L3-L4 for firewall and NAT, L5 for asset monitor and IDS. Besides, NFs actually require more functionality than just packet processing, including file logging and reliable communication with a cluster. File logging needs to deal with file system call. Reliable communication sometimes needs a user-level TCP/IP stack. Correctly integrating these constructs with packet processing pipeline is not a trivial task to do, especially when considering the sequential nature of poll-based packet processing.

We argue that the above mentioned problems can be well handled by proper software abstraction and

the recent advancement in high-performance server side programming.

To effectively integrate a wide range of NFs using a systematic framework, we aim to rely on a widely-known asynchronous programming abstraction called *promise* []. The promise abstraction has the two crucial benefits for implementing future NFs. First, the promise abstraction is expressive enough for implementing a wide range of NFs, from simple packet processing application like firewall and NAT, to complicated NFs that analyze the payload of a flow like asset monitor and IDS. Second, the promise abstraction allows the programmer to write NF code in a synchronous style; the complexity of calling asynchronous callbacks are completely hidden by the promises.

To handle the advanced functionalities that are required by NFs, we leverage the recent advancement in server-side programming, called Seastar []. Seastar is fully compatible with the promise abstraction. It provides high-speed DMA-based file I/O and a user-level TCP/IP stack that is fully compliant with the promise abstraction, making it a perfect base building block for future NF softwares.

In this paper, we propose AsyncNF, which is a framework for building and porting a wide range of NF softwares. AsyncNF provides efficient framework for writing NF software based on promise abstraction. To demonstrate the usability of AsyncNF, we use AsyncNF to implement a wide range of NFs, including a reimplementation of NetBricks, a high-speed port of PRADs asset monitor, a NF replication scheme that does not rely on rolling-back and a simple port of mOS. All the four NF softwares achieve good performance and easy programmability.

The contributions of the paper is listed as follow:

First, we propose a framework for building NF software using promise abstraction. The framework is capable of implementing a wide-range of NFs with advanced functionality such as DMA-based file logging and reli-

able communication.

Second, we perform four case studies using the proposed framework, demonstrating the effectiveness of the proposed framework.

2 Background and Motivation

In recent years, the research community has witnessed the quick development of network function virtualization (NFV). DPDK [2] and Netmap [14] use kernel bypassing to speed up the performance of NF software. They have become the default libraries for implementing high-speed modern NF software. NFV management systems such as E2 [13] are built to dynamically scale virtual instances running different NFs. NFs are augmented with fault tolerance [15] and flow migration [10] to improve the failure resilience.

However, despite all these advancements, a core problem is not well-solved by existing work: what should be the default programming abstraction for implementing NF software, so that the diverse requirements of NF software can be well-captured by this abstraction? To show the importance of this problem, let me first discuss the diversity of NF software.

2.1 Diversity of NF Software

Simple Packet Processing Program. Example NFs include firewall, NAT and load balancer. The word “simple” actually means that the way that these NFs manipulate packets is simple: they take an input packet, perform necessary packet transformation and book-keeping, then they release the packet to the outside. Taking NAT as an example. After receiving an input packet, NAT may update the connection status associated with the flow, then the NAT performs an address translation to substitute the IP address and port of the packet. Finally, NAT sends the packet out from the output port.

These NFs can be effectively implemented inside a polling loop and can be seamlessly integrated with either DPDK or Netmap for maximum performance.

NFs with Intensive File I/O. Example NFs include PRADS [4] asset monitor and Snort [8] intrusion detection system (IDS). For instance, PRADS is a passive real-time asset detection system, which listens to network traffic and logs important information on hosts and services it sees on the network. This information can be used to map the underlying network, letting network operators know what services and hosts are active, and can be used together with IDS/IPS setup for “event to application” correlation.

Both PRADS and Snort can be ported to use DPDK to speed up packet processing [11]. Even after porting to

DPDK, both NFs fail to achieve 10Gbps line rate processing [11]. The primary reason for this undesirable number is due to logging. After porting to DPDK, the worker threads of both NFs keep polling for new packets and maintain CPU usage to 100%. But when both NFs log important events, they have to access system calls related to file system processing, generating expensive context switches and compromising the packet processing throughput.

These NFs can be accelerated using DPDK and Netmap, but they still need to step into the kernel to log events to the files. NFs with intensive file I/O remain to be interesting phenomena in existing NFV research. People have strived to remove context switches associated with kernel networking stack by bypassing the kernel with DPDK, but they fail to remove the context switches associated with kernel file systems during logging.

NFs with Reliable Communication to External Services. Example NFs include S-CSCF in IMS system [1] and NFs that need to replicate their states on back NFs.

S-CSCF is an important middlebox sitting at control plane of the IMS system. It processes SIP [7] messages by contacting several external services. Taking the S-CSCF implementation of a famous open source IMS project Clearwater [5] as an example, when processing SIP messages, S-CSCF needs to log SIP registration information on a Memcached [3] cluster and acquire user information by querying a dedicated storage server called Home Subscriber Server (HSS). The S-CSCF implementation of Clearwater uses kernel TCP/IP stack to carry out reliable communication to all the required external services, seriously limiting the maximum throughput that S-CSCF can achieve. Our experience with Clearwater shows that a single worker thread in S-CSCF can only process SIP messages with the bandwidth of 40Mb.

FTMB [15] is the state of art system for NF replication. It employs a primary-backup replication strategy. On the primary NF instance, after each packet is processed, the packet is passed to the backup over a reliable communication channel for replication. We can treat the replication process as communicating external services: each input packet processed by the primary instance must be reliably delivered to the backup instance. FTMB uses DPDK to speed up packet processing and implements its own reliable communication channel on top of DPDK. But the implementation detail of the reliable communication channel is omitted from the paper. It would be desirable to implement the reliable communication channel using a user-level TCP/IP stack like mTCP [12], so that the performance of FTMB is stable (a handcrafted reliable communication channel may be unstable and lack of flow control) and it is easier to reproduce FTMB implementation for both academic and industrial usage. *However, without a good programming*

abstraction, integrating a user-level TCP implementation like mTCP with replication strategy like FTMB is not a trivial task: mTCP exposes an event-driven programming interface like Linux epoll. The application thread using mTCP does not sit in the same thread as the mTCP worker thread. But FTMB requires that the same worker thread handles both NF packet processing and reliable communication to ensure correct replication.

Some of these NFs abandoned DPDK and Netmap, use kernel networking stack to provide reliable communication channel, but sacrifice performance. Some of these NFs use DPDK and Netmap to speed up packet processing and implement their own reliable communication channel, but sacrifice the stable performance and flow control provided by TCP/IP.

NFs that Process Events Raised by Lower-level System Components. Example NFs include Snort IDS [8] and Bro IDS [9]. The two IDSes alert potential attacks by matching the flow protocols and analyzing flow payloads with an automaton. They can be decoupled into two parts: A low-level system is responsible for re-assembling the TCP stream and generating events associated with the TCP stream, i.e. connection setup, packet re-transmission, and the new packet payload. A high-level event driven system is responsible for reacting to the events raised by the low-level system, i.e. in the case of a fake re-transmission forged by an attacker, the IDS drops the flow and raises an alert. These IDSes can be effectively accelerated using mOS [11], which substitute the low-level system that raises flow-related events. mOS is accelerated using DPDK and is an improved version of mTCP [12].

The low-level system of these NFs can be accelerated with DPDK. However, the low-level system like mOS is usually targeted to process TCP/IP protocol and can not be extended to process non-TCP/IP protocol.

Summary. Now we briefly discuss the similarities and differences of all the discussed NF software.

Similarity. Most of these NFs can be accelerated with DPDK or Netmap (except for S-CSCF, which relies on kernel networking stack, but we can still accelerate it by porting it to user-level TCP/IP stack like mTCP). Using DPDK or Netmap means that the worker threads in these NFs become busy polling thread that keeps CPU usage to 100%, implying that any system calls entering the kernel context may compromise the performance of these NFs.

Difference. These NFs have different working goals and operate at different levels. Simple packet processing programs only manipulate raw packets. They do not rely on any external services. PRADs needs to do file I/O. FTMB and S-CSCF need to communicate with external services. Snort and Bro operate on a high-level that reacts to flow-level events raised by a low-level system components. These differences lead to diverse imple-

mentation details, making it hard to find an appropriate abstraction to unify these NFs.

2.2 One Abstraction to Rule Them All

Just like the dedication that physicists put into the grand unified theory, computer scientists also have been searching for a unified programming abstraction that can capture a variety of applications. In terms of NFV, if a unified programming abstraction can be found for all the NFs mentioned in the previous section, programmers can enjoy the following benefits.

First, by optimizing the performance of the library that provides the unified programming abstraction, we can improve the performance for a huge variety of NFs. There is no need to optimize each NF, which might take a huge amount of labor work.

Second, ease NF software development. Once the implementor becomes familiar with the programming abstraction, he is able to create different types of NFs without learning different programming paradigms or constructing different libraries.

Finally, it makes important research and industrial result easily reproducible, as the unified programming abstraction makes people play on the same ground.

Such a programming abstraction is readily accessible for NFV implementors and researchers, which is functional reactive programming, especially the subset related to futures, promises and continuations.

2.3 Futures and Promises

Futures and promises are important terminologies in functional reactive programming. A future represents a value that is going to be computed while a promise represents the action when the computation is done. This simple programming paradigm can easily capture most of the asynchronous programming patterns. Let me briefly explain how futures and promises can be mapped to NFs discussed in previous sections.

For simple packet processing program, the futures are packets that are going to be received whereas the promises are packet handler functions.

For PRADs, the futures and promises can be combined to implement efficient file system logging. The futures represent the logging action that will log events raised by PRADs to the file system. The promises represents post actions when the logging is done.

For S-CSCF and FTMB, futures and promises can be used to implement an efficient user-space TCP/IP stack. The futures are still packets to be received, but the promises become TCP/IP stack handlers.

For Snort and Bro, futures and promises can be used to implement a low-level system that raises flow events.

Futures flow events that are going to be raised, promises are event handlers for these events.

The future-promise programming abstraction can be efficiently implemented with a small runtime overhead (i.e. asynchronous C++ library Seastar [6]). The programming abstraction can fully bypass the entire kernel, even in terms of file logging (with the help of DMA), providing satisfactory performance for modern NF software.

2.4 Contribution

In this paper, we are going to make the following contributions.

First, we are the first to apply functional reactive programming as a generic method for building a variety of NF software. We use seastar as the underlying library for providing the reactive programming abstraction.

Second, we carry out case studies to show how functional reactive programming can be used to construct 4 different types of NFs, with diverse requirements.

Finally, we show that the performance and ease of implementation are greatly improved by using functional reactive programming. In particular:

We re-implement PRADs using functional reactive programming. The resulting PRADs is capable of logging to file system at a throughput of several gigabits per second.

We create a new primary-backup replication strategy. The new primary-backup strategy is capable of processing packets at line rate. The biggest difference between this replication strategy with FTMB is that it does not need to checkpoint the master NF instance, greatly simplifying the implementation effort. It has no replay time and introduces no extra latency caused by checkpointing.

We re-implement mOS using our new programming abstraction. We also port PRADs to use the new mOS and show that the new PRADs can be several times faster than that in the mOS paper [11].

3 Rollback Recovery for Middleboxes

3.1 FTMB and Non-determinism in Multi-threaded NF

FTMB [15] is the state-of-art middlebox replication strategy. FTMB regularly creates a checkpoint for the primary NF instance and generates a series of packet access logs. In the case of primary NF instance failure, FTMB recover the primary NF instance by rollback: a new primary NF instance is re-created based on the latest checkpoint and the state of the new primary NF instance is recovered by replaying the packet access log.

FTMB simply rejected the idea to keep the primary NF instance and the backup NF instance synchronized, due to the existence of shared variables and non-determinism caused by concurrently accessing shared variables.

A multi-threaded NF software may maintain several shared variables, protected by locks. Accessing these shared variables from multiple worker threads may result in non-determinism: when the same input sequence are fed into two identical programs, they may fail to generate the same output.

For instance, the NF software maintain a shared variable v . There are two NF instances running the same NF software. Both NF instances are configured with two worker threads t_1 and t_2 . Then, two identical packets p_1 and p_2 are concurrently fed into worker threads t_1 and t_2 respectively, on both of the two NF instances. Because the two worker threads run in parallel, it is impossible to guarantee the order for accessing the shared variable v . We might end-up with the following situation: on the first NF instance, t_1 first accesses v when processing p_1 , while on the second NF instance, t_2 first accesses v when processing p_2 . This situation may render the two identical NF instances to be in different states.

Due to the non-determinism, the authors of FTMB [15] reject the design to keep primary and backup synchronized: without tagging the packet processing order from the master instance, it is impossible to keep the backup instance synchronized. Even if master instance correctly tags the packet processing order and sends these tags to the backup instance, the worker threads on backup instance must strictly follow the packet processing order when accessing the shared variable: before accessing a shared variable, the worker thread must keep spinning to wait for its turn to access the shared variable, compromising the performance of the backup instance.

3.2 Execution Model of FTMB

FTMB uses rollback recovery, which could be summarized as follows: the primary instance create a **Packet Access Log**(PAL for short) whenever accessing a shared variable. For each processed packet, the primary instance sends the processed packet, together with all the PALs generated during processing this packet, to the backup instance over a reliable communication channel. The backup instance in FTMB is actually referred to as a **Output Logger**. The output logger stores the PALs for rollback recovery and release the processed packet to the outside. In the meantime, since the primary instance runs in a virtual machine, the primary instance also sends a checkpoint of its virtual machine state to the output logger. The output logger saves the current checkpoint, discards the previous checkpoint and all the PALs received since the previous checkpoint.

If the primary instance crashes, the primary instance is recovered as follow: the output logger uses the saved checkpoint of the primary instance to create a new virtual machine. Then the output logger sends the PALs and the input packets that trigger the generation of these PALs to the new primary instance for replay. The primary instance resumes working after it finishes replaying all the PALs.

3.3 Pros and Cons of FTMB

Pros:

Good system performance. FTMB is carefully designed not to introduce small overhead when generating PALs. The replicated NF in FTMB is able to process millions of packets per second.

Fast recovery time. Depending on how often the primary instance is checkpointed, the recovery time in the case of primary failure can be as small as 20 milliseconds.

Passive Operation. There is no dedicated backup instance in FTMB. The output logger is capable of replicating multiple primary instances.

Cons:

Packet processing latency is increased during checkpoint. When doing checkpoint for primary instance, the primary instance is significantly slowed down, resulting in packet processing latency up to several milliseconds. This leads to a huge jitter in the network RTT, which may be devastating to some important applications such as VoIP and on-line gaming.

The tradeoff between packet latency and checkpoint frequency. An important parameter to tune in FTMB is how frequently should the primary instance be checkpointed. If the primary instance is frequently checkpointed, then the recovery time would be faster, but each time when checkpoint is triggered, FTMB adds huge packet processing latency, sometimes up to several milliseconds. If the primary instance is not frequently checkpointed, then the output logger needs a huge buffer to buffer input packets and PAL, the recovery time is also prolonged.

Lacking An Open-source Solution. FTMB is a very complicated system, involving several critical system components that are non-trivial to engineer. However, FTMB is not open-sourced and therefore limiting its further application in both academia and industry.

4 Middlebox Recovery without Rolling Back

On contrary to FTMB, we aim to provide a new middlebox replication strategy which keeps both the primary and the backup instances in synchronization. This




Figure 1: Work flow of how to keep both primary and backup in synchronization.

replication strategy directly gives the same input packet stream to both the primary instance and the backup instance, our new architecture is able to keep both of the two instances in synchronization.

Using this strategy, there is no need to roll-back the backup instance when the primary instance fails. This strategy minimizes the recovery time and eliminates the prolonged packet processing delay when checkpointing the primary instance.

To tackle the challenge of keeping both the primary instance and the backup instance in synchronization, one might resolve to deterministic scheduling []. However, the overhead caused by deterministic scheduling is way too high for NF software, as a typical NF software needs to process millions of input packets every second. Instead, we solve the deterministic execution problem using a combination of coroutine and message passing.

4.1 Basic Setup for Recovery without Rolling Back

The basic setup is shown in figure 1a. We set up two identical NF instances. The two instances run the same NF software binary image and are configured with the same number of CPU cores.

In 1a, the NF instance is configured with two worker threads (t_1 and t_2) which poll the NIC card for input packet. The shared variable is hosted on a dedicated thread t_3 .

To access the shared variable hosted on t_3 , both t_1 and t_2 need to send a message for accessing the shared variable to t_3 . After the shared variable is modified, another message is sent back to t_1 or t_2 to indicate the completion of shared variable modification.

When the primary instance finishes processing the input packet, it forwards the input packet to the backup instance over a reliable communication channel. The backup instance processes the input packet again before releasing the packet. The packet processed by the backup instance is tagged with an execution order by the primary, so that the shared variable on the backup instance can process the input packet in the same sequence as the primary instance. This guarantees that the state of both

the primary instance and the backup instance are always synchronized.

4.2 Workflow on Primary Instance

Figure 1b shows how worker thread t_2 processes an input packet. The overall workflow is similar to a typical packet polling loop. The only exception is that when a shared variable is going to be accessed by t_2 (step 3 in figure 1b), instead of directly acquiring the lock and update the shared variable, t_2 sends the packet to t_3 to update the shared variable (step 4). When t_3 receives this packet, the threads update the shared variable (step 5), tags the packet with a sequence number (step 6) and sends the packet back to t_2 . When t_2 receives this packet, t_2 sends the tagged packet out to the backup instance over a reliable communication channel.

4.2.1 The Sequence Number

The sequence number tagged by t_3 indicates a sequential accessing order to the shared variable. Using this sequence number, the backup instance can reliably reproduce the accessing order of the shared variable (to be discussed in section 4.3). This ensures that the state of the primary and backup instances are always synchronized.

4.2.2 Using Future and Promise to Cancel Thread Blocking

The biggest problem with this workflow is that, after step 4 in figure 1b, worker thread t_2 must block its execution and wait for the packet to come back from t_3 . This is unacceptable for a high-performance NF software.

We tackle this problem using futures and promises in reactive programming. After step 4 in figure 1b is executed, we wrap the current thread context inside a future object. t_2 can immediately start processing other input packets. When the packet comes back to t_2 after step 6, t_2 is able to re-construct the previous thread context using the corresponding future object. This efficiently eliminates thread blocking.

4.2.3 The Reliable Communication Channel

Due to the power of future and promise, a user-space TCP/IP stack could be integrated inside t_2 . The reliable communication channel is actually a TCP connection channel. This reliable communication channel is augmented with flow control and is more reliable than other specially-crafted reliable communication protocols.

4.3 Workflow on Backup Instance

Figure 1c shows how worker thread t_2 processes the output packet sent from the primary instance. The overall workflow is similar to that of the primary instance. However, when the packet is delivered to t_3 to access the shared variable, t_3 must check whether it has processed all the packets whose sequence number is smaller than the packet. Considering the case of figure 1c, t_3 must wait for packet with sequence number 0 first. If t_3 has processed packet with sequence number 0, t_3 can directly process the packet with sequence number 1. Otherwise, t_3 should store packet with sequence number 1 and wait for the packet with sequence number 0 to come.

Since the order of how packets access the shared variable is well preserved, the backup instance and the primary instance have the same state.

4.4 Recovery

If the primary instance fails, the backup instance can become the primary instance immediately. The recovery time is basically decreased to zero.

5 Promises and Cooperative Threads

In this section, I will first give a detailed overview about promises and cooperative threads used by Seastar framework. Then I will introduce how to apply promises and cooperative threads to NFV.

5.1 Lwt

The core building blocks of Seastar are promises and cooperative threads. Clearly, these fancy concepts come from the world of functional programming languages. There is an Ocaml library called Lwt [16], which also implements promises and cooperative threads. In the following sections, I will first discuss how promises and cooperative threads are implemented in Ocaml. Then I will show how they are implemented in Seastar.

5.1.1 Lwt Overview

When writing programs like network servers, non-blocking is a very important property that ensures the runtime efficiency of the network servers. There are two dominant techniques to make the program non-blocking. The first one uses multi-threading to support non-blocking, i.e. whenever a blocking operation is going to be made, the main thread of the program launches another thread to handle the blocking operation. The second one uses event-based programming method, i.e. the

main thread treats the completion of the blocking operation as an event and register corresponding event handlers to handle this event.

The first technique is easy to use and easy to program. The program can be written in traditional way without relying on callbacks. However, the first technique lacks efficiency, as launching too many threads compromises the runtime performance. On the other hand, the second technique has superior runtime performance, as it only maintains a single thread. However, it is difficult to write programs using the second technique due to excessive use of callbacks.

Lwt uses promises and cooperative threads to support non-blocking. The runtime performance of Lwt is very good, as it only maintains a single physical thread like the second technique. And it is quite easy to use. Writing asynchronous, fully non-blocking programs using Lwt is just like writing synchronous, blocking programs using the first technique.

In the following sections, to facilitate understanding, I simplify some concepts related with exception handling and modify the name of some important types in Lwt.

5.1.2 Promise State

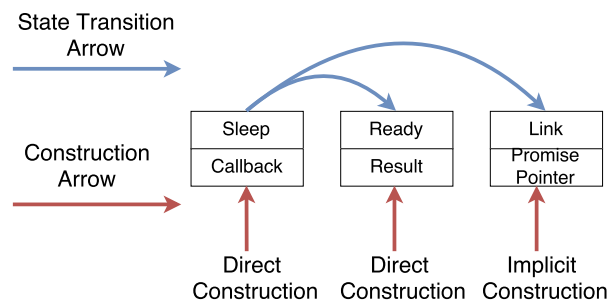


Figure 2: The state of a promise. The blue arrow represents how future states are transitioned. The red arrow represents whether can a programmer construct a certain state.

The basic building block in Lwt is promise. A promise, as shown in figure 2, has three states, which are **Sleep** state, **Ready** state and **Link** state. The **Sleep** state represents that the result of the promise is not available yet and one has to wait before the **Sleep** state transitions into **Ready** state. The **Sleep** state promise may contain a callback function, which is immediately called after it is transitioned to **Ready** state. **Ready** state represents that the result is available and we can peek the result by checking the result field. Finally, the **Link** state contains a pointer to a promise that is in **Sleep** state. It is used to when waiting for multiple events.

The state of a promise can be changed. The blue arrow of figure 2 shows the state transition graph between

the three states. Only **Sleep** state can generate a state transition.

When programming with promises, the programmer can only construct **Sleep** state promise and **Ready** state promise. The **Link** state promise is implicated constructed when chaining a promise with an anonymous function (we omit the discussion, as it does not affect the understanding of how promises work).

5.1.3 Chaining Promises

`(>>=) : Promise a -> (a -> Promise b) -> Promise b`

Figure 3: The infix operator for chaining promises.

Multiple promises can be chained together to accomplish complicated tasks. Chaining is done through an infix operator `>>=` (the then member function of future in Seastar) as shown in figure 3, whose signature is listed below.

Here, Promise a represents a promise that, when transitioning into **Ready** state, contains a result field with type a. `(a -> Promise b)` represents an anonymous function, that takes a value of type a as argument and returns a value with type Promise b. `>>=` is a function, that takes a value of Promise a and an anonymous function of `(a -> Promise b)` and returns a value of Promise b.

The real power of the `>>=` operator is to chain multiple promises together into a complicated operation. We give a piece of example code in figure 4. The code will first sleep for 3 seconds, then print "first print" on the screen, then sleep another 3 seconds, and finally print "second print" on the screen. The best thing about this code is that, even if it represents the consecutive execution of four blocking operations, but the code itself is not blocking at all. It will return a promise in **Sleep** state after being called. The blocking operations are implicitly handled by a background thread.

```
sleep 3 >>=
fun () -> async_print "first print" >>=
fun () -> sleep 3 >>=
fun () -> async_print "second print"
```

Figure 4: Chaining multiple promises into a complicated operation. The code above will first sleep for 3 seconds, then print "first print" on the screen, then sleep another 3 seconds, and finally print "second print" on the screen.

5.1.4 Anatomy of the infix operator

The internal of the infix operator. Depending on the state of promise x, the infix operator may step into two

```

(>>=) x f =
  match x with
  | Result r -> f r
  | Sleep ->
    let res = make_sleep_promise () in
    add_callback x ( fun x -> connect res (x >>= f) )
    res
  | Link p ->
    assert false
connect t t' =
  match t' with
  | Result r -> (run t's callback)
  | Sleep ->
    (let t' link to t)

```

Figure 5: The implementation of the infix operator.

branches. If x is in **Result** state, then the anonymous function is immediately executed. But if x is in **Sleep** state, the operator first creates a new promise res in **Sleep** state. Then a callback is added to x . When x becomes **Result** state, the callback function is called, which connects res and $x \gg f$. The `connect` function means that the state of promise t' will be reflected in promise t .

A conceptual explanation. When executing $x \gg= f$, if x is immediately available and in **Result** state, then the execution of the anonymous function f continues without any interruption.

The tricky part comes when x is in **Sleep** state. It implies that the result of x is going to be available in the future. To prevent the infix operator from blocking, we construct a new **Sleep** state promise res and return res to the user. Apparently, res should reflect the actual result of $x \gg= f$. To achieve this, before returning res to the user, we add a callback function to x , which manually connect res with $x \gg= f$ when x becomes ready. In this way, we manually construct a fully non-blocking abstraction.

5.1.5 Wake Up Sleep Promise

If a promise is in **Sleep** state, it must be waken up in the future and transform into **Result** state. This is done by an external polling thread which polls asynchronous completion messages. If the blocking operation that a **Sleep** promise waits for has completed, the corresponding promise is waken up.

5.2 Seastar

Seastar is basically a re-implementation of `lwt` in C++. It differs from `lwt`:

First, the infix operator $\gg=$ that chains promises together becomes `then` member function.

Second, Seastar introduces a new object called `future`, which is actually a pointer object to the underlying

promise. The reason is that C++ has no garbage collection and Seastar has to manually manage the memory allocation of the promise object. Currently, Seastar either allocates the promise object directly on the heap, or captures the promise object inside the callback function. Therefore, to query the promise, seastar creates a pointer object `future`, which contains a pointer to the underlying promise object.

5.3 Promises in NFV

We use promises to hide any kind of blocking operations when processing packets. We give some usage examples.

5.3.1 Simple Packet Processing

```

xxx.then([]{
  process packet;
  return new_future;
}).then([]{
  process packet;
  return new_future;
})

```

Figure 6: Example code of a simple packet processing software.

It is straightforward to implement a simple packet processing software using promises, as shown in figure 6. The multiple chained promises represent multiple processing stages on a packet processing pipeline.

5.3.2 Using Promises to Hide Blocking During File IO

```

xxx.then([]{
  do file IO;
  return new_future;
}).then([]{
  resume packet processing;
  return new_future;
})

```

Figure 7: Example code of a NF that performs file IO.

Figure 7 represents a NF that performs file IO in the middle of packet processing. Previously, this incurs kernel context switches and blocking, which may compromise the performance of the NF. But with promise, the kernel context switches and blocking are hidden by the promises, making the entire operation fully non-blocking.

After issuing `do file IO`, the NF software can continue to process other flows. When file IO finishes, the processing of the original packet that incur the file IO is

resumed. This can greatly boost the performance of NFs that need to perform file IO, such as PRADs.

5.3.3 Handling Shared Variable Accessing

```
xxx.then([]{
    access shared variable;
    return new_future;
}).then([]{
    resume packet processing;
    return new_future;
})
```

Figure 8: Example code of a NF that needs to access shared variable.

Figure 8 shows an example of how to use promise to handle shared variable processing. After issuing `access shared variable`, the processing is temporarily suspended and a request is created and sent to another thread to access the shared variable. After modifying the shared variable, the suspended packet processing is resumed.

By using promises, we can effectively linearize shared variable accessing. Therefore we can keep the primary NF instance and the backup NF instance in the same state, which simplifies primary-backup replication even in multi-threaded environment.

References

- [1] 3GPP specification: 23.228. <http://www.3gpp.org/ftp/Specs/html-info/23228.htm>.
- [2] Intel Data Plane Development Tool Kit. <http://dpdk.org/>.
- [3] Memcached. <https://memcached.org/>.
- [4] PRADs Asset Monitor. <https://github.com/gamelinux/prads>.
- [5] Project Clearwater. <http://www.projectclearwater.org/>.
- [6] Seastar Project. <http://www.seastar-project.org/>.
- [7] SIP: Session Initiation Protocol. <https://www.ietf.org/rfc/rfc3261.txt>.
- [8] Snort intrusion detection system. <https://www.snort.org/>.
- [9] The Bro Network Security Monitor. <https://www.bro.org/>.
- [10] GEMBER-JACOBSON, A., VISWANATHAN, R., PRAKASH, C., GRANDL, R., KHALID, J., DAS, S., AND AKELLA, A. Opennf: Enabling innovation in network function control. In *Proc. of ACM SIGCOMM* (2014).
- [11] JAMSHED, M. A., MOON, Y., KIM, D., HAN, D., AND PARK, K. mos: A reusable networking stack for flow monitoring middleboxes. In *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)* (Boston, MA, 2017), USENIX Association, pp. 113–129.
- [12] JEONG, E., WOOD, S., JAMSHED, M., JEONG, H., IHM, S., HAN, D., AND PARK, K. mtcp: a highly scalable user-level tcp stack for multicore systems. In *11th USENIX Symposium on Networked Systems Design and Implementation (NSDI 14)* (Seattle, WA, 2014), USENIX Association, pp. 489–502.
- [13] PALKAR, S., LAN, C., HAN, S., JANG, K., PANDA, A., RATNASAMY, S., RIZZO, L., AND SHENKER, S. E2: a framework for nfv applications. In *Proceedings of the 25th Symposium on Operating Systems Principles* (2015), ACM, pp. 121–136.
- [14] RIZZO, L. Netmap: a novel framework for fast packet i/o. In *21st USENIX Security Symposium (USENIX Security 12)* (2012), pp. 101–112.
- [15] SHERRY, J., GAO, P. X., BASU, S., PANDA, A., KRISHNAMURTHY, A., MACIOCCO, C., MANESH, M., MARTINS, J., RATNASAMY, S., RIZZO, L., ET AL. Rollback-recovery for middleboxes. In *ACM SIGCOMM Computer Communication Review* (2015), vol. 45, ACM, pp. 227–240.
- [16] VOULLON, J. Lwt: a cooperative thread library. In *Proceedings of the 2008 ACM SIGPLAN workshop on ML* (2008), ACM, pp. 3–12.